

Repair Service with(in) C*.


Vinay Chella
Senior Software Engineer

NETFLIX

Agenda.

- Challenges in keeping the data consistent
- State of Repairs Today
- Ideal world for repair
- Repair Service @ Netflix
- Roadmap

State of **Repairs in C*** community
these days.



Challenges in keeping data consistent.

- Multi region deployments
- PB's of data
- Types of repairs
 - *Incremental repairs?*
 - Subrange repairs
- Low consistency writes on high throughput clusters
- Flaky networks/ flaky hardware



Today's state of Repairs

- Not easy to setup and orchestrate
 - Reaper(<http://cassandra-reaper.io/>)
 - Tickler
 - Custom repair orchestration
- Stuck repairs
- Failed repairs
- Long running repairs

Ideal world for Repairs.

- Easy to setup
 - Something that Cassandra dev/ admin should not worry about
- Natural as part of C* architecture
 - Forget about repair, yet your data is consistent, but tune it if needed
- Does not put too much burden on the system
- Resilient for failures (Hardware, network and data)



Repair Service @ Netflix

- Decentralized (distributed), self-schedulable, self-healable and self-monitoring service to keep the data consistent across nodes and data centers
- Flaky HW/ flaky network does not impact Repair Service.
- No extra setup needed, Always ON by default, C* tries to keep your data consistent, can be tuned using table properties
- Not to put extra load on C*, yet repair intelligently
- Tune GC_GRACE_SECONDS based on repair durations when needed
- Monitor repair thread pools to avoid stuck repairs

Future of this project

- Data aware repair / infectious repair
- Adaptive scheduling, load for repair
- Drifted range repairs
- Integrating as part of C* @ Netflix
- As part of Priam(C* sidecar) now, C* would be the ideal home for it
- Similar JIRA - 10070



Repair Service



cassandra




cassandra

Repair Service

Repair Service *within* C*.

Repairs

C*



Take away

Let's change C^* from *hopefully* consistent to *eventually* consistent



Thank you.

NETFLIX

